

CSCE 435 Fall 2015

HW 2: Parallel Programming with Threads II

Due: 11:59pm Friday, September 25, 2015

Minimum of a List

You are provided with a program `list_minimum.c` to compute the minimum element of a list using threads. The main program initializes data structures and calls `pthread_create` to create threads that execute the function `find_minimum`. Each thread is assigned a sublist ranging from `my_start` to `my_end` and is responsible for computing the minimum value in this range. The thread is also responsible for updating the global variable `minimum` to ensure that the global minimum is no larger than the minimum value in its sublist. **The program is missing code to update `minimum`.** You need to add code to the `find_minimum` routine so that the program computes the minimum correctly. Access to `minimum` is controlled by the mutex `lock_minimum`. The variable `count` has been initialized to zero in the main program, and can be used to determine how many threads have accessed `minimum`.

Once you complete the code, it can be compiled with the command:

```
icc -o list_minimum.exe list_minimum.c -lpthread -lc
```

Note the use of `-lc` to link the library that provides the random number generator. To execute the program, use

```
./list_minimum.exe <n> <p>
```

where `<n>` represents the number of elements in the list and `<p>` represents the number of threads. The output of a sample run is shown below.

```
./list_minimum.exe 100000 8
```

```
Threads = 8, minimum = 18074, time (sec) = 0.0006
```

1. (20 points) Complete the function `find_minimum` so that the program computes the minimum of the list correctly. You will get full points if you add code only to the `find_minimum` routine.
2. (10 points) Execute the code for $n = 2 \times 10^8$ with p chosen to be 2^k , for $k = 0, \dots, 13$. Plot execution time versus p to demonstrate how time varies with the number of threads. Use logarithmic scale for the x-axis. Plot speedup versus p to demonstrate the change in speedup with p .
3. (10 points) Give reasons for the observed variation in execution time as p is varied from 1 to 8192.

Barrier

You are provided with a program `barrier.c` that has code to implement a barrier for a multithreaded program. The main program initializes data structures and calls `pthread_create` to create threads that execute the function `work` before calling the barrier function. The file `csce435.h` includes `work` that forces a thread to sleep for a specified time before returning. **The program is missing code for the barrier function called `barrier_simple`.** You need to add code to `barrier_simple` to implement a barrier among the threads. Threads should wait for all threads to enter the barrier routine before they are allowed to exit. The variable `count` can be used

to determine how many threads have entered the barrier. You may use the mutex `lock_barrier` to allow a thread to obtain exclusive access to `count`. The condition variable `cond_barrier` can be used to have threads wait to receive a signal when `count` reaches a specific value.

4. (20 points) Complete the function `barrier_simple` to implement a barrier among the threads. You will get full points if you add code only to the `barrier_simple` routine.
5. (10 points) Execute the code for $p=2^k$, for $k = 1, \dots, 14$. Plot execution time versus p to demonstrate how time varies with the number of threads. Use logarithmic scale for the x-axis.
6. (10 points) Set `sleeptime.tv_sec=0` and `sleeptime.tv_nsec=0` in `csce435.h` and run the experiments from the previous step again. Plot execution time versus p . How would you characterize the growth in time with p ? What is the reason for such a growth?

List Statistics

7. (20 points) Modify the program in `list_minimum.c` so that it computes the mean and standard deviation of the list elements instead of the minimum. Name the new program file `list_statistics.c`. You may define global variables `mean` and `standard_deviation` that will store the values. You will get full credit only if the mean and standard deviation of the list are computed by threads before exiting the thread routine.

Submission: You need to upload the following to ecampus:

1. Problem 1: submit the file `list_minimum.c`.
2. Problem 4: submit the file `barrier.c`; you do not need to submit `csce435.h`.
3. Problem 7: submit the file `list_statistics.c`.
4. Problems 2, 3, 5, and 6: submit a single PDF or MSWord document with your responses.

Helpful Information:

1. Source files are available on eCampus.
2. Load the compiler module prior to compiling your program. Use:
`module load intel`
3. Compile C programs using `icc`. Link the `pthread` library when using `pthread`. For example, to compile `code.c` to create the executable `code.exe`, use
`icc -o code.exe code.c -lpthread`
4. You may need to use `dos2unix` on the files downloaded from `ecampus.tamu.edu` to strip additional characters that may get introduced by eCampus.
5. The run time of a code should be measured when it is executed in dedicated mode. Create a batch file, say `batch_file`, as described on `sc.tamu.edu` and use the following command to submit it to the batch system:
`bsub < batch_file`