

## CSCE 435 Fall 2015

### HW 7: Graph Algorithms

Due: 11:59pm Monday, December 7, 2015

Given a fully connected undirected weighted graph  $G = (V, E, w)$  with **positive** edge weights, the minimum spanning tree (MST) is defined as a spanning tree with minimum weight. A spanning tree of an undirected graph is a tree containing all nodes of the graph. The weight of a spanning tree is the sum of weights of all the edges in the tree.

You are provided a file called `mst.c` that has the following routines:

- `init_adj_matrix`: generates the adjacency matrix  $A$  of an undirected weighted graph where the weights between edges are selected randomly;
- `minimum_spanning_tree`: computes the MST for  $A$  using Prim's algorithm.

The code uses two data structures:

- `ADJ_MATRIX_t`: stores the adjacency matrix as a two-dimensional array. `A[i][j]` stores the weight of the edge between nodes  $i$  and  $j$ . The number of nodes in the graph is stored in `n`.
- `MST_t`: stores the MST in the form of two arrays called `node` and `weight`.

Prim's algorithm grows the MST one node at a time starting from a root node. A node  $u$  is included in the current MST if the weight of the edge from  $u$  to an MST node, say  $v$ , is the smallest weight edge among all edges between MST nodes and non-MST nodes. The array element `node[u]` stores  $v$  and `weight[u]` stores the weight of the edge  $(u, v)$ . Thus, the spanning tree is represented by the pairs  $(i, \text{node}[i])$  for  $i = 0, \dots, n-1$  ( $i \neq \text{root}$ ), and the corresponding edge weights are stored in `weight[i]`.

1. (80 points) You need to develop a parallel implementation of the minimum spanning tree algorithm that computes the MST of  $A$ . You are allowed to select one of the following approaches to parallelize the code:
  - a) pthreads-based implementation on ADA that uses all the cores on a single ADA node; or
  - b) CUDA-based implementation on ADA that uses a GPU device.

A total of 20 points are reserved for performance of the code: speedup obtained by the parallel code over the serial code in `mst.c`.

2. (20 points) Execute the code for various values of  $n$  to demonstrate the parallel performance of your code. Include material such as plots and arguments as appropriate to convince the reader that your implementation delivers high parallel performance.

**Submission:** You need to upload the following to ecampus:

1. Problem 1: submit the file `mst.c`.
2. Problem 2: submit a single PDF or MSWord document with your response.